You already told us what you didn’t understand.

Echo — English learners hit rewind dozens of times a day, and no app records that tap.
We catch it: no questions, no interruption. Tomorrow you get those fifteen seconds back.

Personalisation in language apps is fake.

IS THE PAIN REAL?

They teach **their syllabus**, not the English you actually miss. I have studied English for over a decade and still lose half of a native podcast — every non-native speaker here knows the feeling.

Today's answer — Anki, a vocabulary notebook — asks you to **flag it yourself**. That action is a tax, which is why almost everyone quits after two days.

HOW BIG IS IT?

About **1.5 billion** people are learning English. The ones who hurt most, and pay most readily, are those **already living in English** and still not catching all of it — international students, expat workers, engineers on remote teams.

I am that person. Six weeks in the Bay Area, hit by this every day.

Every rewind is **an honest confession**.
People already produce this signal every day — nobody is catching it.

ZERO FRICTION

They never decide “I want to learn this”. **The rewind was going to happen anyway.**

NO INTERRUPTION

Nothing is asked while they listen. Every judgment call is **deferred to the next day**.

DATA, NOT A FEATURE

It grows **a personal map of your English** — a rival can copy the feature, not this.

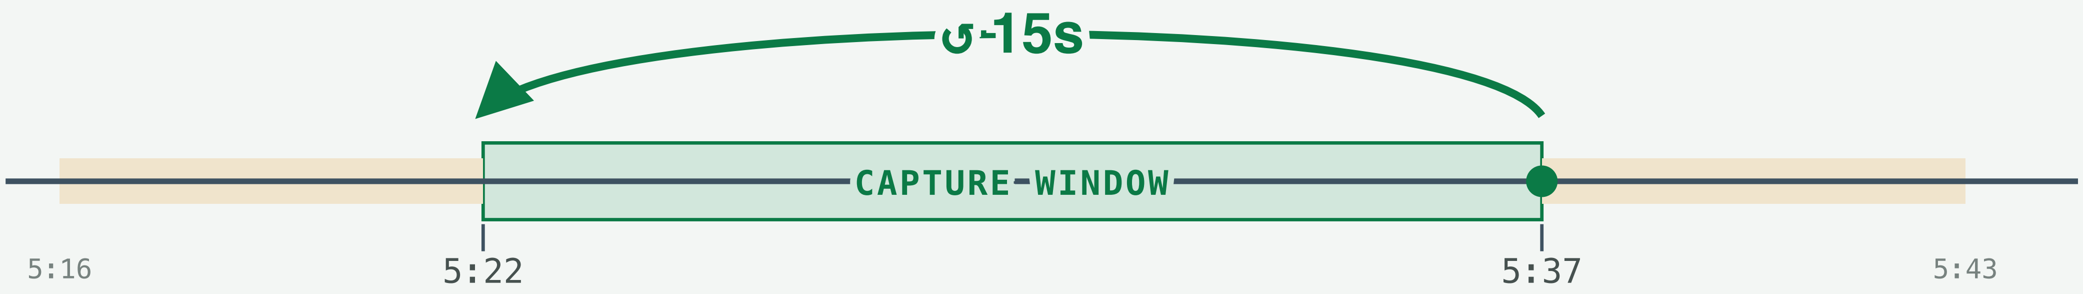

AS IT HAPPENS	SECONDS LATER	THAT EVENING	NEXT MORNING	AFTER THAT
Catch the tap Back button, Control Center, lock screen all count	**Transcribe just this** Windowed ASR, not the whole episode	**Diagnose the break** Vocabulary / linking / speed / grammar / accent / culture	**Reveal in stages** Clue → skeleton → full text; the mask keeps the sentence's shape	**Ask again before they forget** Spaced repetition

From the rewind to the question next morning, all of it is automatic. They did nothing extra.

Watch. I hit rewind.

1. Listening to a podcast, one sentence doesn't land → **tap ↺15**. Nothing happens on screen
2. Switch to Practice → ⚡ **Get ahead (1)**: those fifteen seconds walked in by themselves
3. The sentence is **masked** — but the mask keeps its shape: length and rhythm are still there
4. Listen first, then decide: **"I really didn't get it" or "I was just distracted"**
5. Three-stage reveal: clue → skeleton → full text

WHY THE CLUE IS THE *PREVIOUS* SENTENCE

Listening runs forward. **The previous sentence is the one they already understood** — that is what makes it a foothold. The next one would just be the ending.

WHY THE MASK HAS A SHAPE

They can see how long the sentence is, which word is unusually long, what the rhythm is — but they cannot read it. **Shape is exactly the cue listening uses to rebuild a sentence.**

WHY NOW

- ASR cost fell an order of magnitude in two years; **windowed transcription** brings the unit cost into range
- LLMs only just became able to answer “**why didn’t they understand it**”, not merely “what was said”
- Bluetooth earbuds = **a billion sensors already deployed**, with programmable gestures
- The ambient-assistant wave on AR glasses is just starting — **same signal, new carrier**

MOAT: NOT THE FEATURE, THE DATA

All three pillars have been built by someone. **The moat is at the input layer** — capturing difficulty from real listening. Nobody is collecting that.

> Features can be copied. **The map of your English that three months of rewinds builds cannot.**

And to copy it you must first **rewrite the player** — only your own player sees that event. That is why we are not a web app.

TAM

$15–20B

Global **digital language learning** (industry estimates). ~1.5 billion English learners worldwide.

SAM

$3–4B

The **English × listening/speaking-led** self-study slice. English ≈ 50% × listening-and-speaking-led ≈ 40%.

SOM (3 YEARS)

$16M ARR

A **bottom-up** beachhead. Every assumption is spelled out below.

HOW THE SOM IS BUILT — EVERY ASSUMPTION HERE IS CHALLENGEABLE, AND THAT IS THE POINT

Non-native students and early-career professionals in the US **≈ 1.5 M** (~1.1 M international students + work-visa stock) × already listen to English podcasts **30%** = 450 K × willing to pay for English learning **20%** = 90 K × ARPU **$15/mo** (against Speak / ELSA at $15–25) = **≈ $16M ARR**

The beachhead is here because these people **hurt the most, pay the most readily, and I am standing among them**.

WHO	WHAT THEY HAVE	WHAT THEY LACK
Snipd closest	Triple-tap earbud capture of podcast clips, ASR + LLM summaries	Serves knowledge management; **no language diagnosis or pronunciation layer at all**
Duolingo Video Call	Live AI speaking practice and feedback, at incumbent scale	Syllabus-driven; **no passive listening signal**
Speak $1B valuation	Voice-first AI tutor; proved the high-ARPU subscription	You practise its prompts, **not what you actually missed**
YourBestAccent et al.	Clone your own voice to practise pronunciation (golden speaker)	No real-life input, **no diagnostic loop**
Meta Ray-Ban	Wearable live translation; hardware already shipping at scale	Does **translation**, not **learning** — the biggest risk to Phase 3
Echo	**Rewind signal from real listening → diagnosis of why → practice loop**	**All three exist separately; nobody has connected them**

⚠️ I do **not** claim voice cloning or AI speaking diagnosis as mine — both are red oceans, and the evidence for self-voice pronunciation practice is mixed. **I am betting on one thing only: the input layer.**

HOW IT MAKES MONEY

Subscription, $15/month. Free: unlimited capture and transcripts. Paid: diagnosis, practice cards, spaced repetition, shadowing in your own voice.

The cost story is **windowed transcription** — transcribe the fifteen seconds they replayed, not the episode. Annotating a 10-minute window costs about **$0.005**.

> **But talking about pricing now is empty.** I have not measured retention. What has to be proven first is the **next-day return rate**: will they open it tomorrow to practise what they rewound today?

BUILT IN 5 WEEKS (ALL CHECKABLE)

- **End-to-end loop running**: rewind → transcript → diagnosis → practice → schedule, fully automatic
- **Lock-screen rewinds are caught too**: 3 measured, jumps of 10.0–10.1 s, matching the system's hardcoded interval
- **27 architecture decision records** — every trade-off with its reason and its cost written down
- 14 dependencies; a one-line JS change **reaches the phone in 60 seconds**, no re-review

Real users: 1 (me, dogfooding, week 5). A working pipeline is not a useful product — I know the difference, which is why the next slide is what I want tonight.

TEAM: RIGHT NOW, ME

I am the target user — non-native, in the Bay Area, hit by this every day. NTU CSIE; I ran an N=12 within-subject study on self-voice shadowing, so I know **which claims hold up and which don't**.

What I'm missing is specific: a speech/ML engineering co-founder, and a first cohort who will actually use it every day.

THREE THINGS I WANT TONIGHT

- **Ten people who'll use it seven days straight.** Do you learn English from podcasts? You can install it right now
- **A co-founder with a speech/ML background** — or an introduction to one
- **One introduction to an early-stage investor or accelerator**

Booth 08. Come hit the rewind button on my phone — thirty seconds later you'll watch it turn into data.