



You already told us what you didn't **understand.**

Echo — English learners hit rewind dozens of times a day, and no app records that tap.

We catch it: no questions, no interruption. Tomorrow you get those fifteen seconds back.



Personalisation in language apps is **fake**.

IS THE PAIN REAL?

They teach **their syllabus**, not the English you actually miss. I have studied English for over a decade and still lose half of a native podcast — every non-native speaker here knows the feeling.

Today's answer — Anki, a vocabulary notebook — asks you to **flag it yourself**. That action is a tax, which is why almost everyone quits after two days.

HOW BIG IS IT?

About **1.5 billion** people are learning English. The ones who hurt most, and pay most readily, are those **already living in English** and still not catching all of it — international students, expat workers, engineers on remote teams.

I am that person. Six weeks in the Bay Area, hit by this every day.



Every rewind is **an honest confession**.
People already produce this signal every day — nobody is catching it.

ZERO FRICTION


They never decide “I want to learn this”. **The rewind was going to happen anyway.**

NO INTERRUPTION

Nothing is asked while they listen. Every judgment call is **deferred to the next day.**

DATA, NOT A FEATURE

It grows **a personal map of your English** — a rival can copy the feature, not this.



AS IT HAPPENS

SECONDS LATER

THAT EVENING

NEXT MORNING

AFTER THAT

Catch the tap

Back button, Control
Center, lock screen all
count

Transcribe just this

Windowed ASR, not the
whole episode

Diagnose the break

Vocabulary / linking /
speed / grammar / accent
/ culture

Reveal in stages

Clue → skeleton → full text;
the mask keeps the
sentence's shape

Ask again before they forget

Spaced repetition

From the rewind to the question next morning, all of it is automatic. They did nothing extra.

Watch. I hit **rewind**.

- 1 Listening to a podcast, one sentence doesn't land → **tap** ⏮15. Nothing happens on screen
- 2 Switch to Practice → ⚡ **Get ahead (1)**: those fifteen seconds walked in by themselves
- 3 The sentence is **masked** — but the mask keeps its shape: length and rhythm are still there
- 4 Listen first, then decide: **"I really didn't get it" or "I was just distracted"**
- 5 Three-stage reveal: clue → skeleton → full text

WHY THE CLUE IS THE *PREVIOUS* SENTENCE

Listening runs forward. **The previous sentence is the one they already understood** — that is what makes it a foothold. The next one would just be the ending.

WHY THE MASK HAS A SHAPE

They can see how long the sentence is, which word is unusually long, what the rhythm is — but they cannot read it. **Shape is exactly the cue listening uses to rebuild a sentence.**

WHY NOW

- ASR cost fell an order of magnitude in two years; **windowed transcription** brings the unit cost into range
- LLMs only just became able to answer “**why didn’t they understand it**”, not merely “what was said”
- Bluetooth earbuds = **a billion sensors already deployed**, with programmable gestures
- The ambient-assistant wave on AR glasses is just starting — **same signal, new carrier**

MOAT: NOT THE FEATURE, THE DATA

All three pillars have been built by someone. **The moat is at the input layer** — capturing difficulty from real listening. Nobody is collecting that.

Features can be copied. **The map of your English that three months of rewinds builds cannot.**

And that input layer **keeps widening**: your own podcasts today, then YouTube and Netflix, ending at earbuds capturing everyday talk. **The closer to real life, the better the signal.**

And to copy it you must first **rewrite the player** — only your own player sees that event.

TAM

\$15–20B

Global **digital language learning** (industry estimates).
~1.5 billion English learners worldwide.

SAM

\$3–4B

The **English × listening/speaking-led** self-study slice.
English ≈ 50% × listening-and-speaking-led ≈ 40%.

SOM (3 YEARS)

\$16M ARR

A **bottom-up** beachhead. Every assumption is spelled out below.

HOW THE SOM IS BUILT — EVERY ASSUMPTION HERE IS CHALLENGEABLE, AND THAT IS THE POINT

Non-native students and early-career professionals in the US ≈ **1.5 M** (~1.1 M international students + work-visa stock) × already listen to English podcasts **30%** = 450 K × willing to pay for English learning **20%** = 90 K × ARPU **\$15/mo** (against Speak / ELSA at \$15–25) = ≈ **\$16M ARR**

The beachhead is here because these people **hurt the most, pay the most readily, and I am standing among them.**

WHO	WHAT THEY HAVE	WHAT THEY LACK
Snipd closest	Triple-tap earbud capture of podcast clips, ASR + LLM summaries	Serves knowledge management; no language diagnosis or pronunciation layer at all
Duolingo Video Call	Live AI speaking practice and feedback, at incumbent scale	Syllabus-driven; no passive listening signal
Speak \$1B valuation	Voice-first AI tutor; proved the high-ARPU subscription	You practise its prompts, not what you actually missed
YourBestAccent et al.	Clone your own voice to practise pronunciation (golden speaker)	No real-life input, no diagnostic loop
Meta Ray-Ban	Wearable live translation; hardware already shipping at scale	Does translation , not learning — the biggest risk to Phase 3
Echo	Rewind signal from real listening → diagnosis of why → practice loop	All three exist separately; nobody has connected them

⚠️ I do **not** claim voice cloning or AI speaking diagnosis as mine — both are red oceans, and the evidence for self-voice practice is mixed (my own N=12: subjective preference swept all four dimensions, **objective measures did not separate**). But I have **already shipped that chain** — Mirror · AI Speech Coach is live on the App Store, 502 signups in 4 weeks, zero ad spend. It is a capability I hold, not a risk I carry. **I am still betting on one thing only: the input layer.**

HOW IT MAKES MONEY

Subscription, \$15/month. Free: unlimited capture and transcripts. Paid: diagnosis, practice cards, spaced repetition, shadowing in your own voice.

The cost story is **windowed transcription** — transcribe the fifteen seconds they replayed, not the episode. Annotating a 10-minute window costs about **\$0.005**.

But talking about pricing now is empty. I have not measured retention. What has to be proven first is the **next-day return rate**: will they open it tomorrow to practise what they rewound today?

BUILT IN 5 WEEKS (ALL CHECKABLE)

- **End-to-end loop running**: rewind → transcript → diagnosis → practice → schedule, fully automatic
- **Lock-screen rewinds are caught too**: 3 measured, jumps of 10.0–10.1 s, matching the system's hardcoded interval
- **27 architecture decision records** — every trade-off with its reason and its cost written down
- 14 dependencies; a one-line JS change **reaches the phone in 60 seconds**, no re-review

Real users: 1 (me, dogfooding, week 5). A working pipeline is not a useful product — I know the difference, which is why the next slide is what I want tonight.

TEAM: RIGHT NOW, ME

I am the target user — non-native, in the Bay Area, hit by this every day. NTU CSIE. My last product, **Mirror · AI Speech Coach**, is live on the **App Store**: 502 signups in 4 weeks, zero ad spend, 90.2% onboarding completion. Its N=12 study is why I know **which claims hold up and which don't**.

What I'm missing is specific: a speech/ML engineering co-founder, and a first cohort who will actually use it every day.

THREE THINGS I WANT TONIGHT

- **Ten people who'll use it seven days straight.** Do you learn English from podcasts? You can install it right now
- **A co-founder with a speech/ML background** — or an introduction to one
- **One introduction to an early-stage investor or accelerator**

Booth 08. Come hit the rewind button on my phone — thirty seconds later you'll watch it turn into data.