


You already told us what you didn't understand.



Language learners hit rewind dozens of times a day, and **no app records that tap**. Echo keeps the fifteen seconds you replayed — no questions, no interruption — and hands them back tomorrow.



01 / AS IT HAPPENS

Catch the tap

Back button, Control Center, lock screen. Taps within 3 s count once.

02 / SECONDS LATER

Transcribe just this

Windowed ASR — not the whole episode. Cheaper, and no waiting.

03 / THAT EVENING

Diagnose the break

Vocabulary, linking, speed, grammar, accent, culture.

04 / NEXT MORNING

It comes to you

A notification carries the word to the lock screen. Three choices, plus “can’t recall”.

05 / AFTER THAT

Before you forget

Spaced repetition; every grade leaves a review log.

Why the rewind button is a good signal

3.06
ODDS
RATIO

Learners who **spontaneously** noticed a vocabulary hole were 3.06× as likely to learn that word.

4.57
ODDS
RATIO

Learners **prompted** to speak, and who noticed the gap that way: 4.57×.

n.s.
BETWEEN
THE TWO

The two groups **did not differ significantly** — forcing the action adds nothing over noticing it yourself.

What does the work is **realising you have a gap** — not the action. Hitting rewind *is* that moment. So we never interrupt.

Bilingualism: Lang. & Cognition · Field (2003), Lexical Segmentation in L2 Listening

The difference is the input

Duolingo-style apps feed you a **syllabus** — words and patterns you never missed. Echo’s input is the sentence you didn’t catch today, from a source you pick.

your podcasts → YouTube · Netflix → earbuds, real talk

The moat is not any one feature. It is that input plus what sits behind it: AI difficulty diagnosis, label analysis, own-voice shadowing — none of it English-specific.

I have shipped the own-voice piece before. Mirror · AI Speech Coach is live on the App Store — 502 signups in 4 weeks, zero ad spend, 90.2% onboarding. But my N=12 study is precise: **subjective preference swept all four dimensions; objective measures did not separate**. It buys willingness to speak, not efficacy — Echo still grades on the original.

Working today

- ✓ **End-to-end loop** — rewind to scheduled practice
- ✓ **Lock-screen rewinds count too** — pressed with the screen off
- ✓ **Scheduled lock-screen prompts** — the word is the notification title; a wrong answer reschedules it for today
- ✓ **Two-tap selection**, plus an “I can’t tell where the words split” escape hatch
- ✓ **Difficulty bands 1–5** — measured and guessed look different



Try it — or tell me where I’m wrong

9 questions · about 90 seconds. One TestFlight invite when the beta opens. No newsletter.

rick-ray-wldd.github.io/daily-learn-app